

Data
Visualization
and Pre-
processing

Description:- Predicting the age of abalone from physical measurements. The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-onsuming task. Other measurements, which are easier to obtain, are used to predict age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem.

Building a Regression Model

Perform Below Visualizations.

Univariate Analysis

Summary Statistics

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from google.colab import files
```

```
uploaded = files.upload()
```

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving abalone.csv to abalone.csv

```
file_data = pd.read_csv('abalone.csv')
```

```
file_data
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Rings	
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10

```

4174    M      0.600  0.475  0.205  1.1760  0.5255  0.2875  0.3080  9
4175    F      0.625  0.485  0.150  1.0945  0.5310  0.2610  0.2960 10
4176    M      0.710  0.555  0.195  1.9485  0.9455  0.3765  0.4950 12
4177 rows x 9 columns

```

Add a Age column in a dataset

```
file_data['Age']=''
```

```
file_data.head()
```

```

Sex      Length Diameter      Height Whole weight  Shucked weight Viscera
weight Shell weight  Rings  Age
0      M      0.455  0.365  0.095  0.5140  0.2245  0.1010  0.150  15
1      M      0.350  0.265  0.090  0.2255  0.0995  0.0485  0.070  7
2      F      0.530  0.420  0.135  0.6770  0.2565  0.1415  0.210  9
3      M      0.440  0.365  0.125  0.5160  0.2155  0.1140  0.155  10
4      I      0.330  0.255  0.080  0.2050  0.0895  0.0395  0.055  7

```

```
file_data['Age']=file_data['Rings']+1.5
```

```
file_data.head()
```

```

Sex      Length Diameter      Height Whole weight  Shucked weight Viscera
weight Shell weight  Rings  Age
0      M      0.455  0.365  0.095  0.5140  0.2245  0.1010  0.150  15
      16.5
1      M      0.350  0.265  0.090  0.2255  0.0995  0.0485  0.070  7      8.5
2      F      0.530  0.420  0.135  0.6770  0.2565  0.1415  0.210  9
      10.5
3      M      0.440  0.365  0.125  0.5160  0.2155  0.1140  0.155  10
      11.5
4      I      0.330  0.255  0.080  0.2050  0.0895  0.0395  0.055  7      8.5

```

Drop the Rings Column

```
file_data = file_data.drop(columns=['Rings'],axis=1)
```

```
file_data
```

```

Sex      Length Diameter      Height Whole weight  Shucked weight Viscera
weight Shell weight  Age
0      M      0.455  0.365  0.095  0.5140  0.2245  0.1010  0.1500 16.5
1      M      0.350  0.265  0.090  0.2255  0.0995  0.0485  0.0700 8.5
2      F      0.530  0.420  0.135  0.6770  0.2565  0.1415  0.2100 10.5
3      M      0.440  0.365  0.125  0.5160  0.2155  0.1140  0.1550 11.5
4      I      0.330  0.255  0.080  0.2050  0.0895  0.0395  0.0550 8.5
...      ...      ...      ...      ...      ...      ...      ...
4172    F      0.565  0.450  0.165  0.8870  0.3700  0.2390  0.2490 12.5
4173    M      0.590  0.440  0.135  0.9660  0.4390  0.2145  0.2605 11.5
4174    M      0.600  0.475  0.205  1.1760  0.5255  0.2875  0.3080 10.5

```

```

4175    F      0.625  0.485  0.150  1.0945  0.5310  0.2610  0.2960  11.5
4176    M      0.710  0.555  0.195  1.9485  0.9455  0.3765  0.4950  13.5
4177 rows x 9 columns

```

```
file_data['Height'].mean()
```

```
0.13951639932966242
```

```
file_data['Height'].median()
```

```
0.14
```

```
file_data['Height'].std()
```

```
0.041827056607257274
```

```
Frequency Table
```

```
file_data['Sex'].value_counts()
```

```
M    1528
```

```
I    1342
```

```
F    1307
```

```
Name: Sex, dtype: int64
```

```
Create Charts
```

```
file_data.boxplot(column=['Length'], grid=False)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f04b69ef390>
```

```
file_data.hist(column='Length', grid=False, edgecolor='black')
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f04b695c910>]],
      dtype=object)
```

```
sns.kdeplot(file_data['Length'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f04b6430a50>
```

```
Bi - Variate Analysis
```

```
Barplot
```

```
data = sns.barplot(x = file_data["Sex"], y = file_data["Age"])
```

```
data
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f04b6404450>
```

```
Correlation Coefficients
```

```
file_data.corr()
```

```

Length Diameter      Height Whole weight  Shucked weight Viscera weight
      Shell weight  Age
Length  1.000000      0.986812      0.827554      0.925261      0.897914
      0.903018      0.897706      0.556720
Diameter      0.986812      1.000000      0.833684      0.925452
      0.893162      0.899724      0.905330      0.574660

```

Height	0.827554	0.833684	1.000000	0.819221	0.774972
	0.798319	0.817338	0.557467		
Whole weight	0.925261	0.925452	0.819221	1.000000	
	0.969405	0.966375	0.955355	0.540390	
Shucked weight	0.897914	0.893162	0.774972	0.969405	
	1.000000	0.931961	0.882617	0.420884	
Viscera weight	0.903018	0.899724	0.798319	0.966375	
	0.931961	1.000000	0.907656	0.503819	
Shell weight	0.897706	0.905330	0.817338	0.955355	
	0.882617	0.907656	1.000000	0.627574	
Age	0.556720	0.574660	0.557467	0.540390	0.420884
	0.503819	0.627574	1.000000		

3.Linear Plot

```
data = sns.lineplot(x = file_data["Length"], y = file_data["Age"])
data
<matplotlib.axes._subplots.AxesSubplot at 0x7f04b697b410>
```

Scatter Plot

```
data = sns.scatterplot(x = file_data['Age'],y = file_data['Diameter'],
color="green")
data
<matplotlib.axes._subplots.AxesSubplot at 0x7f04b6319390>
```

Multi - Variate Analysis

```
x = sns.scatterplot(x=file_data['Whole
weight'],y=file_data['Age'],hue=file_data["Sex"])
x
<matplotlib.axes._subplots.AxesSubplot at 0x7f04b6234910>
```

Perform descriptive statistics on the dataset.

```
file_data.shape
(4177, 9)
file_data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0    Sex              4177 non-null   object
1    Length           4177 non-null   float64
2    Diameter         4177 non-null   float64
```

3	Height	4177 non-null	float64
4	Whole weight	4177 non-null	float64
5	Shucked weight	4177 non-null	float64
6	Viscera weight	4177 non-null	float64
7	Shell weight	4177 non-null	float64
8	Age	4177 non-null	float64

dtypes: float64(8), object(1)

memory usage: 293.8+ KB

file_data.describe()

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
		Shell weight	Age			
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	
	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	
mean	0.523992	0.407881	0.139516	0.828742	0.359367	
	0.180594	0.238831	11.433684			
std	0.120093	0.099240	0.041827	0.490389	0.221963	
	0.109614	0.139203	3.224169			
min	0.075000	0.055000	0.000000	0.002000	0.001000	
	0.000500	0.001500	2.500000			
25%	0.450000	0.350000	0.115000	0.441500	0.186000	
	0.093500	0.130000	9.500000			
50%	0.545000	0.425000	0.140000	0.799500	0.336000	
	0.171000	0.234000	10.500000			
75%	0.615000	0.480000	0.165000	1.153000	0.502000	
	0.253000	0.329000	12.500000			
max	0.815000	0.650000	1.130000	2.825500	1.488000	
	0.760000	1.005000	30.500000			

file_data.head()

Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
		Shell weight	Age			
0	M	0.455	0.365	0.095	0.5140	0.2245
1	M	0.350	0.265	0.090	0.2255	0.0995
2	F	0.530	0.420	0.135	0.6770	0.2565
3	M	0.440	0.365	0.125	0.5160	0.2155
4	I	0.330	0.255	0.080	0.2050	0.0895

file_data.tail()

Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
		Shell weight	Age			
4172	F	0.565	0.450	0.165	0.8870	0.3700
4173	M	0.590	0.440	0.135	0.9660	0.4390
4174	M	0.600	0.475	0.205	1.1760	0.5255
4175	F	0.625	0.485	0.150	1.0945	0.5310
4176	M	0.710	0.555	0.195	1.9485	0.9455

```
file_data.mean(numeric_only=True)
```

```
Length          0.523992
Diameter        0.407881
Height          0.139516
Whole weight    0.828742
Shucked weight  0.359367
Viscera weight  0.180594
Shell weight    0.238831
Age             11.433684
```

```
dtype: float64
```

```
file_data.median(numeric_only=True)
```

```
Length          0.5450
Diameter        0.4250
Height          0.1400
Whole weight    0.7995
Shucked weight  0.3360
Viscera weight  0.1710
Shell weight    0.2340
Age             10.5000
```

```
dtype: float64
```

```
file_data.mode()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
0	M	0.550	0.45	0.15	0.2225	0.175	0.1715	0.275	10.5
1	NaN	0.625	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
file_data.var(numeric_only=True)
```

```
Length          0.014422
Diameter        0.009849
Height          0.001750
Whole weight    0.240481
Shucked weight  0.049268
Viscera weight  0.012015
Shell weight    0.019377
Age             10.395266
```

```
dtype: float64
```

```
file_data.std(numeric_only=True)
```

```
Length          0.120093
Diameter        0.099240
Height          0.041827
Whole weight    0.490389
Shucked weight  0.221963
Viscera weight  0.109614
Shell weight    0.139203
```


4175	False	False	False	False	False	False
4176	False	False	False	False	False	False

	Viscera weight	Shell weight	Age
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
...
4172	False	False	False
4173	False	False	False
4174	False	False	False
4175	False	False	False
4176	False	False	False

```
[4177 rows x 9 columns]
print(file_data.isnull().sum())
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
Age          0
dtype: int64
file_data.isna().any()
Sex          False
Length       False
Diameter     False
Height       False
Whole weight False
Shucked weight False
Viscera weight False
Shell weight False
Age          False
dtype: bool
Find the outliers and replace the outliers
x = sns.boxplot(x=file_data["Age"])
x
<matplotlib.axes._subplots.AxesSubplot at 0x7f04b6152ad0>
```



```

6      0      0      0      0      0      0      0      0      0      0      0
    ...      0      0      0      0      0      0      0      0      0      0
    0
7      0      0      0      0      0      0      0      0      0      0      0
    ...      0      0      0      0      0      0      0      0      0      0
    0
8      0      0      0      0      0      0      0      0      0      0      0
    ...      0      0      0      0      0      0      0      0      0      0
    0
9      0      0      0      0      0      0      0      0      0      0      0
    ...      0      0      0      0      0      0      0      0      0      0
    0

```

10 rows x 51 columns

```
pd.get_dummies(file_data).head(10)
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight				
		Shell weight	Age	Sex_F	Sex_I	Sex_M				
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5	0	0
	1									
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5	0	0
	1									
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5	1	0
	0									
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5	0	0
	1									
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5	0	1
	0									
5	0.425	0.300	0.095	0.3515	0.1410	0.0775	0.120	9.5	0	1
	0									
6	0.530	0.415	0.150	0.7775	0.2370	0.1415	0.330	21.5	1	0
	0									
7	0.545	0.425	0.125	0.7680	0.2940	0.1495	0.260	17.5	1	0
	0									
8	0.475	0.370	0.125	0.5095	0.2165	0.1125	0.165	10.5	0	0
	1									
9	0.550	0.440	0.150	0.8945	0.3145	0.1510	0.320	20.5	1	0
	0									

Split the data into dependent and independent variables.

Splitting the Dataset into the Independent

```
X = file_data.iloc[:, :-1].values
```

```
print(X)
```

```
['M' 0.455 0.365 ... 0.2245 0.101 0.15]
```

```

['M' 0.35 0.265 ... 0.0995 0.0485 0.07]
['F' 0.53 0.42 ... 0.2565 0.1415 0.21]
...
['M' 0.6 0.475 ... 0.5255 0.2875 0.308]
['F' 0.625 0.485 ... 0.531 0.261 0.296]
['M' 0.71 0.555 ... 0.9455 0.3765 0.495]]
# Extracting the Dataset to Get the Dependent

Y = file_data.iloc[:, -1].values
print(Y)
[16.5  8.5 10.5 ... 10.5 11.5 13.5]
Scale the independent variables
from sklearn.preprocessing import scale
x = scale(file_data["Viscera weight"])
x
array([-0.72621157, -1.20522124, -0.35668983, ...,  0.97541324,
        0.73362741,  1.78744868])
Split the data into training and testing
from sklearn.model_selection import train_test_split
x = file_data.iloc[:, 1:7]
x
Length Diameter      Height Whole weight  Shucked weight Viscera weight
0      0.455  0.365  0.095  0.5140  0.2245  0.1010
1      0.350  0.265  0.090  0.2255  0.0995  0.0485
2      0.530  0.420  0.135  0.6770  0.2565  0.1415
3      0.440  0.365  0.125  0.5160  0.2155  0.1140
4      0.330  0.255  0.080  0.2050  0.0895  0.0395
...      ...      ...      ...      ...      ...
4172  0.565  0.450  0.165  0.8870  0.3700  0.2390
4173  0.590  0.440  0.135  0.9660  0.4390  0.2145
4174  0.600  0.475  0.205  1.1760  0.5255  0.2875
4175  0.625  0.485  0.150  1.0945  0.5310  0.2610
4176  0.710  0.555  0.195  1.9485  0.9455  0.3765
4177 rows x 6 columns

y = file_data.iloc[:, -1]
y
0      16.5
1       8.5
2     10.5
3     11.5
4       8.5
...

```

```
4172    12.5
4173    11.5
4174    10.5
4175    11.5
4176    13.5
```

```
Name: Age, Length: 4177, dtype: float64
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state =42)
```

```
x_train
```

Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	
3823	0.615	0.455	0.135	1.0590	0.4735	0.2630
3956	0.515	0.395	0.140	0.6860	0.2810	0.1255
3623	0.660	0.530	0.175	1.5830	0.7395	0.3505
0	0.455	0.365	0.095	0.5140	0.2245	0.1010
2183	0.495	0.400	0.155	0.8085	0.2345	0.1155
...
3444	0.490	0.400	0.115	0.5690	0.2560	0.1325
466	0.670	0.550	0.190	1.3905	0.5425	0.3035
3092	0.510	0.395	0.125	0.5805	0.2440	0.1335
3772	0.575	0.465	0.120	1.0535	0.5160	0.2185
860	0.595	0.475	0.160	1.1405	0.5470	0.2310

```
3132 rows x 6 columns
```

```
y_train
```

```
3823    10.5
3956    13.5
3623    11.5
0       16.5
2183     7.5
...
3444    10.5
466     13.5
3092    12.5
3772    10.5
860     7.5
```

```
Name: Age, Length: 3132, dtype: float64
```

```
print(x_train.shape, x_test.shape)
```

```
(3132, 6) (1045, 6)
```

```
Build the Model
```

```
from sklearn.linear_model import LinearRegression
```

```
model=LinearRegression()
```

```
model.fit(x_train,y_train)
```

```
LinearRegression()
```

12. Train the Model

```
Y_predict_train = model.predict(x_train)
Y_predict_train
array([11.25888828, 11.95379472, 12.33692259, ..., 11.12903068,
       10.71152746, 11.59516371])
```

13. Test the Model

```
y_predict = model.predict(x_test)
y_predict
array([13.0478407 , 11.43166184, 15.59825921, ..., 13.69440346,
       11.79279231, 10.83037939])
```

Measure the performance using Metrics

```
from sklearn.metrics import mean_squared_error
import math
print(mean_squared_error(y_test, y_predict))
print(math.sqrt(mean_squared_error(y_test, y_predict)))
4.862459933051861
2.2050986220692854
```